

UAS Sistem Terdistribusi – Pub-Sub Log Aggregator

Python 3.11

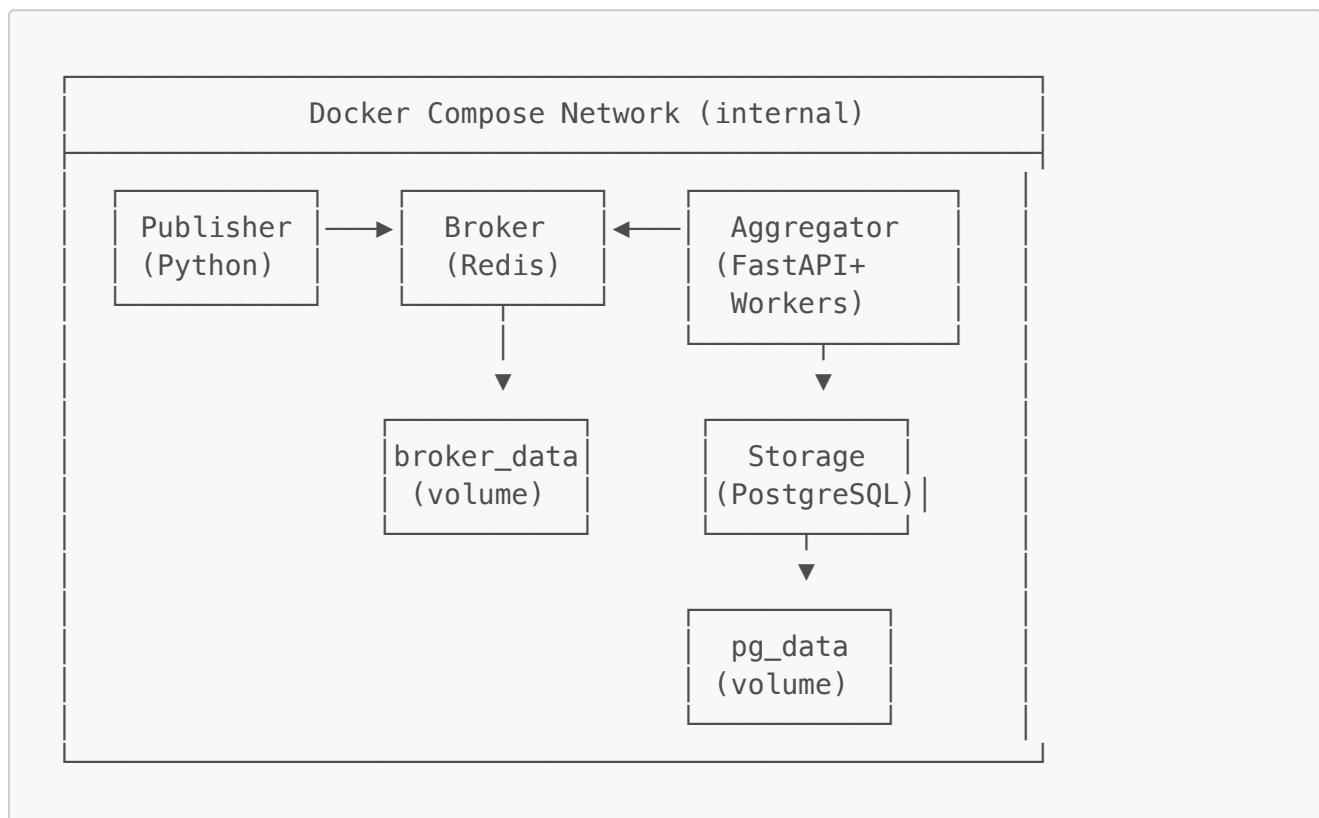
FastAPI 0.110

Docker Compose

Tests 20 passing

Sistem Pub-Sub Log Aggregator terdistribusi dengan **idempotent consumer**, **deduplication kuat**, dan **kontrol transaksi/konkurensi** berbasis Docker Compose.

Arsitektur



Komponen

- **aggregator**: FastAPI + background workers untuk processing events
- **publisher**: Generator event dengan duplikasi terkontrol (30%+)
- **broker**: Redis 7 sebagai message queue internal
- **storage**: PostgreSQL 16 untuk persistent dedup store

Quick Start

Menjalankan Sistem

```
# Build dan jalankan semua services
docker compose up --build
```

```
# Jalankan dengan publisher demo
docker compose --profile demo up --build

# Atau jalankan di background
docker compose up --build -d
```

Akses Endpoints

Endpoint	Method	Deskripsi
http://localhost:8080/health	GET	Health check
http://localhost:8080/publish	POST	Publish events
http://localhost:8080/events?topic=X	GET	List events by topic
http://localhost:8080/stats	GET	Aggregator statistics
http://localhost:8080/queue/stats	GET	Queue status

Contoh Request

```
# Health check
curl http://localhost:8080/health

# Publish single event
curl -X POST http://localhost:8080/publish \
-H "Content-Type: application/json" \
-d '{
  "events": [
    {
      "topic": "user-events",
      "event_id": "evt-12345",
      "timestamp": "2024-01-01T12:00:00Z",
      "source": "demo",
      "payload": {"action": "login"}
    }
  ]
}'

# Get statistics
curl http://localhost:8080/stats

# List events by topic
curl "http://localhost:8080/events?topic=user-events"
```

Testing

Unit/Integration Tests (20 tests)

```
# Install dependencies
pip install -r requirements-dev.txt

# Run all tests
pytest tests/ -v

# Run with coverage
pytest tests/ -v --cov=aggregator
```

Test Coverage

- Deduplication tests (3 tests)
- Persistence tests (2 tests)
- Transaction/Concurrency tests (4 tests)
- Schema validation tests (4 tests)
- API endpoint tests (4 tests)
- Batch processing tests (2 tests)
- Stress/Performance test (1 test)

Load Testing (K6)

```
# Install K6: https://k6.io/docs/getting-started/installation/
# macOS: brew install k6

# Run load test
k6 run loadtest/k6-script.js

# Custom options
k6 run --vus 20 --duration 60s loadtest/k6-script.js
```

Features

Idempotency & Deduplication

- Unique constraint (`topic, event_id`) di PostgreSQL
- `INSERT ON CONFLICT DO NOTHING` untuk atomic dedup
- Metrics: `unique_processed, duplicate_dropped`

Transaksi & Konkurensi

- ACID compliance dengan PostgreSQL
- Isolation level: READ COMMITTED
- Atomic counter updates: `SET count = count + 1`
- Thread-safe workers dengan independent sessions

Reliability

- At-least-once delivery dengan retry eksponensial
- Crash tolerance via persistent storage
- Container restart policy: `unless-stopped`
- Health checks untuk dependency ordering

Persistensi

- Named volumes: `pg_data`, `broker_data`, `aggregator_data`
- Data aman meski container dihapus
- Redis AOF untuk queue durability

Observability

- `/health` endpoint untuk liveness/readiness
- `/stats` dengan dedup rate percentage
- Structured logging
- `/queue/stats` untuk monitoring

Struktur Proyek

```

UAS/
├── aggregator/
│   ├── Dockerfile
│   ├── requirements.txt
│   └── app/
│       ├── main.py          # FastAPI endpoints
│       ├── consumer.py      # Background workers
│       ├── models.py        # SQLAlchemy models
│       ├── schemas.py       # Pydantic schemas
│       ├── db.py            # Database connection
│       ├── queue.py         # Redis/InMemory queue
│       └── config.py        # Configuration
├── publisher/
│   ├── Dockerfile
│   ├── requirements.txt
│   └── app/
│       └── main.py          # Event generator
└── tests/
    └── test_api.py        # 20 unit/integration tests
└── loadtest/
    └── k6-script.js       # K6 load test script
└── scripts/
    ├── demo.sh
    ├── quick-test.sh
    └── test-docker.sh
└── docker-compose.yml    # Main compose file
└── docker-compose.test.yml # Test compose file
└── requirements-dev.txt  # Development dependencies
└── REPORT.md             # Full report (T1-T10)
└── README.md              # This file

```

Configuration

Environment Variables

Variable	Default	Description
DATABASE_URL	postgresql+psycopg2://user:pass@storage:5432/uasdb	Database connection
REDIS_URL	redis://broker:6379/0	Redis connection
WORKER_COUNT	3	Number of background workers
QUEUE_KEY	event_queue	Redis queue key
DISABLE_WORKERS	0	Disable workers (for testing)
USE_INMEMORY_QUEUE	0	Use in-memory queue (for testing)

Publisher Configuration

Variable	Default	Description
TARGET_URL	http://aggregator:8080/publish	Aggregator URL
TOPIC	demo-topic	Default topic
RATE_PER_SEC	100	Events per second
DUP_RATE	0.35	Duplicate rate (35%)

API Schema

Event JSON

```
{  
  "topic": "string (required, 1-255 chars)",  
  "event_id": "string (required, unique per topic)",  
  "timestamp": "ISO8601 datetime",  
  "source": "string (required)",  
}
```

```
"payload": { "any": "json object" }
```

Stats Response

```
{  
  "received": 1000,  
  "unique_processed": 700,  
  "duplicate_dropped": 300,  
  "dedup_rate_percent": 30.0,  
  "topics": ["topic-a", "topic-b"],  
  "uptime_seconds": 120.5  
}
```

Video Demo

Link: [YouTube - UAS Sistem Terdistribusi Demo](#)

References

Tanenbaum, A. S., & Van Steen, M. (2017). *Distributed Systems: Principles and Paradigms* (3rd ed.). Pearson Education.

License

MIT License - Free for educational use.